

dysi5mhay

July 28, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df=pd.read_csv("/content/13_placement.csv")
df
```

```
[2]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
..
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

[1000 rows x 3 columns]

```
[3]: df.head()
```

```
[3]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0

1 DATA CLEANING AND DATA PREPROCESSING

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cgpa                   1000 non-null   float64
1   placement_exam_marks  1000 non-null   float64
2   placed                 1000 non-null   int64
dtypes: float64(2), int64(1)
memory usage: 23.6 KB

```

```
[5]: df.describe()
```

```

[5]:          cgpa  placement_exam_marks  placed
count  1000.000000      1000.000000  1000.000000
mean     6.961240        32.225000    0.489000
std     0.615898        19.130822    0.500129
min     4.890000         0.000000    0.000000
25%     6.550000        17.000000    0.000000
50%     6.960000        28.000000    0.000000
75%     7.370000        44.000000    1.000000
max     9.120000       100.000000    1.000000

```

```
[6]: df.columns
```

```
[6]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

```
[7]: df1=df.dropna(axis=1)
df1
```

```

[7]:          cgpa  placement_exam_marks  placed
0     7.19         26.0         1
1     7.46         38.0         1
2     7.54         40.0         1
3     6.42          8.0         1
4     7.23         17.0         0
..    ...         ...         ...
995   8.87         44.0         1
996   9.12         65.0         1
997   4.89         34.0         0
998   8.62         46.0         1
999   4.90         10.0         1

```

```
[1000 rows x 3 columns]
```

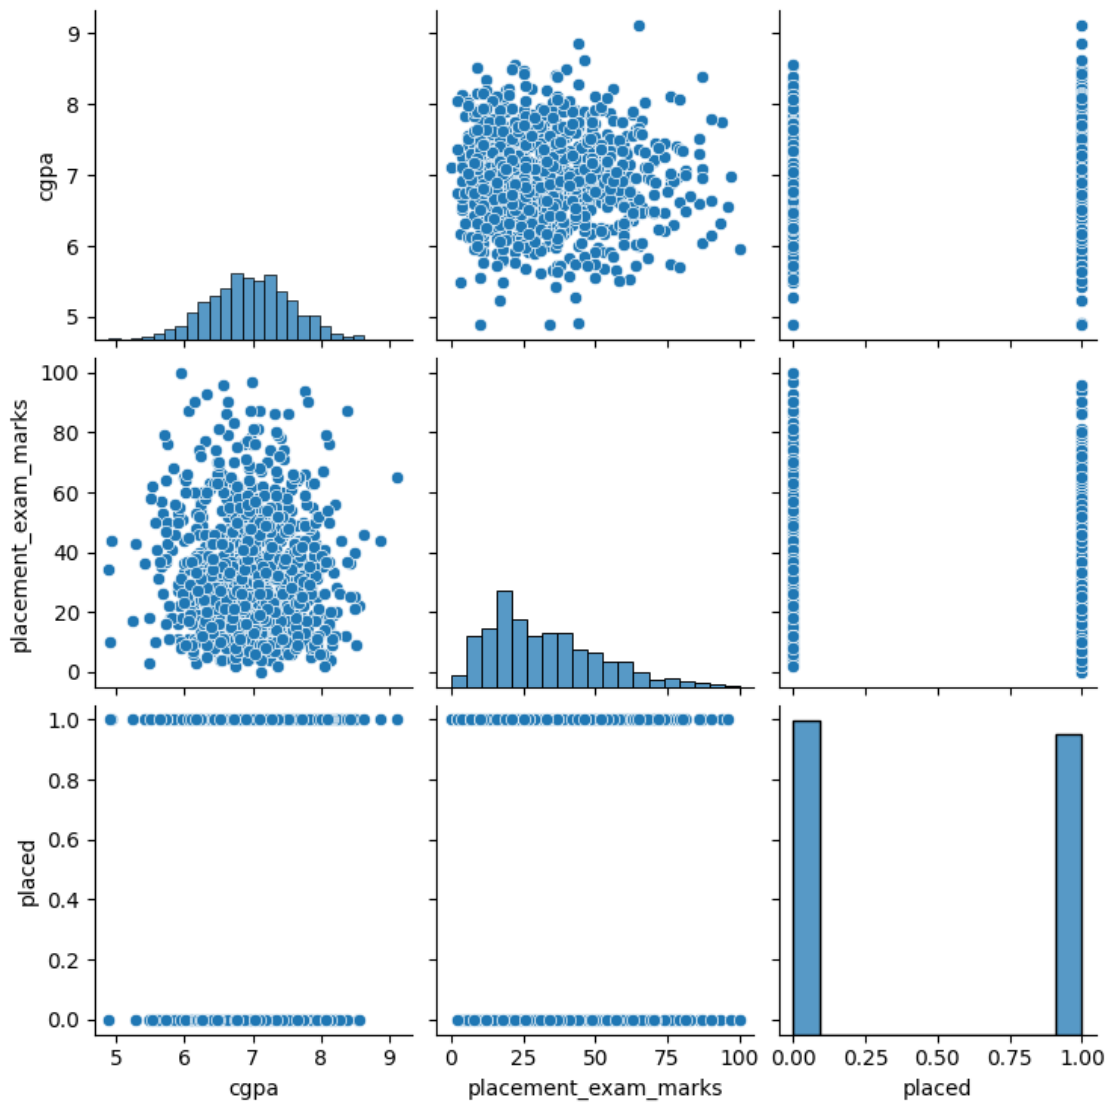
```
[8]: df1.columns
```

```
[8]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

2 EDA AND VISUALIZATION

```
[9]: sns.pairplot(df1)
```

```
[9]: <seaborn.axisgrid.PairGrid at 0x7bbcd99e2020>
```



```
[10]: sns.distplot(df1['placed'])
```

<ipython-input-10-dc9f78aae914>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df1['placed'])
```

```
[10]: <Axes: xlabel='placed', ylabel='Density'>
```



```
[11]: sns.heatmap(df1.corr())
```

```
[11]: <Axes: >
```



3 TO TRAIN THE MODEL AND MODEL BUILDING

```
[12]: x=df[['cgpa', 'placement_exam_marks']]
      y=df['placed']
```

```
[13]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
[14]: from sklearn.linear_model import LinearRegression
      lr=LinearRegression()
      lr.fit(x_train,y_train)
```

```
[14]: LinearRegression()
```

```
[15]: lr.intercept_
```

```
[15]: 0.2577539814877772
```


```
[16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
      coeff
```

```
[16]:
```

	Co-efficient
cgpa	0.035498
placement_exam_marks	-0.000742

```
[17]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
[17]: <matplotlib.collections.PathCollection at 0x7bbed2804a90>
```



4 ACCURACY

```
[18]: lr.score(x_test,y_test)
```

```
[18]: -0.00561591284914309
```

```
[19]: lr.score(x_train,y_train)
```

```
[19]: 0.002765450941136449
```

```
[20]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
```

```
rr.fit(x_train,y_train)
```

[20]: Ridge(alpha=10)

```
[21]: rr.score(x_train,y_train)
```

[21]: 0.0027629711257934897

```
[22]: rr.score(x_test,y_test)
```

[22]: -0.005435920006991557

```
[23]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

[23]: Lasso(alpha=10)

```
[24]: la.score(x_train,y_train)
```

[24]: 0.0

```
[25]: la.score(x_test,y_test)
```

[25]: -0.0025482988358924707