

wbtliapnf

July 28, 2023

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[4]: df=pd.read_csv("/content/1_fiat500_VehicleSelection_Dataset.csv")
df
```

```
[4]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1.0	lounge	51.0	882.0	25000.0	1.0	
1	2.0	pop	51.0	1186.0	32500.0	1.0	
2	3.0	sport	74.0	4658.0	142228.0	1.0	
3	4.0	lounge	51.0	2739.0	160000.0	1.0	
4	5.0	pop	73.0	3074.0	106880.0	1.0	
...
1544	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1545	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1546	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1547	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1548	NaN	NaN	NaN	NaN	NaN	NaN	NaN

	lat	lon	price	Unnamed: 9	Unnamed: 10
0	44.907242	8.611559868	8900	NaN	NaN
1	45.666359	12.24188995	8800	NaN	NaN
2	45.503300	11.41784	4200	NaN	NaN
3	40.633171	17.63460922	6000	NaN	NaN
4	41.903221	12.49565029	5700	NaN	NaN
...
1544	NaN	length	5	NaN	NaN
1545	NaN	concat	lonprice	NaN	NaN
1546	NaN	Null values	NO	NaN	NaN
1547	NaN	find	1	NaN	NaN
1548	NaN	search	1	NaN	NaN

[1549 rows x 11 columns]

```
[5]: df.head()
```

```
[5]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1.0	lounge	51.0	882.0	25000.0	1.0	
1	2.0	pop	51.0	1186.0	32500.0	1.0	
2	3.0	sport	74.0	4658.0	142228.0	1.0	
3	4.0	lounge	51.0	2739.0	160000.0	1.0	
4	5.0	pop	73.0	3074.0	106880.0	1.0	

	lat	lon	price	Unnamed: 9	Unnamed: 10
0	44.907242	8.611559868	8900	NaN	NaN
1	45.666359	12.24188995	8800	NaN	NaN
2	45.503300	11.41784	4200	NaN	NaN
3	40.633171	17.63460922	6000	NaN	NaN
4	41.903221	12.49565029	5700	NaN	NaN

1 DATA CLEANING AND DATA PREPROCESSING

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1549 entries, 0 to 1548
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   float64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   float64
3   age_in_days           1538 non-null   float64
4   km                    1538 non-null   float64
5   previous_owners       1538 non-null   float64
6   lat                   1538 non-null   float64
7   lon                   1549 non-null   object
8   price                 1549 non-null   object
9   Unnamed: 9            0 non-null      float64
10  Unnamed: 10           1 non-null      object
dtypes: float64(7), object(4)
memory usage: 133.2+ KB
```

```
[7]: df.describe()
```

```
[7]:
```

	ID	engine_power	age_in_days	km	previous_owners	\
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	
min	1.000000	51.000000	366.000000	1232.000000	1.000000	
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	

75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000

	lat	Unnamed: 9
count	1538.000000	0.0
mean	43.541361	NaN
std	2.133518	NaN
min	36.855839	NaN
25%	41.802990	NaN
50%	44.394096	NaN
75%	45.467960	NaN
max	46.795612	NaN

```
[8]: df.columns
```

```
[8]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
          'lat', 'lon', 'price', 'Unnamed: 9', 'Unnamed: 10'],
          dtype='object')
```

```
[9]: df1=df[0:1500]
```

```
[10]: df1=df1.dropna(axis=1)
df1
```

```
[10]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1.0	lounge	51.0	882.0	25000.0	1.0	
1	2.0	pop	51.0	1186.0	32500.0	1.0	
2	3.0	sport	74.0	4658.0	142228.0	1.0	
3	4.0	lounge	51.0	2739.0	160000.0	1.0	
4	5.0	pop	73.0	3074.0	106880.0	1.0	
...	
1495	1496.0	pop	62.0	3347.0	80000.0	3.0	
1496	1497.0	pop	51.0	1461.0	91055.0	3.0	
1497	1498.0	lounge	51.0	397.0	15840.0	3.0	
1498	1499.0	sport	51.0	1400.0	60000.0	1.0	
1499	1500.0	pop	51.0	1066.0	53100.0	1.0	

	lat	lon	price
0	44.907242	8.611559868	8900
1	45.666359	12.24188995	8800
2	45.503300	11.41784	4200
3	40.633171	17.63460922	6000
4	41.903221	12.49565029	5700
...
1495	44.283878	11.88813972	7900
1496	44.508839	11.46907997	7450
1497	38.122070	13.36112022	10700

```
1498  45.802021  9.187789917  10800
1499  38.122070  13.36112022   8900
```

```
[1500 rows x 9 columns]
```

```
[11]: df1.columns
```

```
[11]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
         'lat', 'lon', 'price'],
        dtype='object')
```

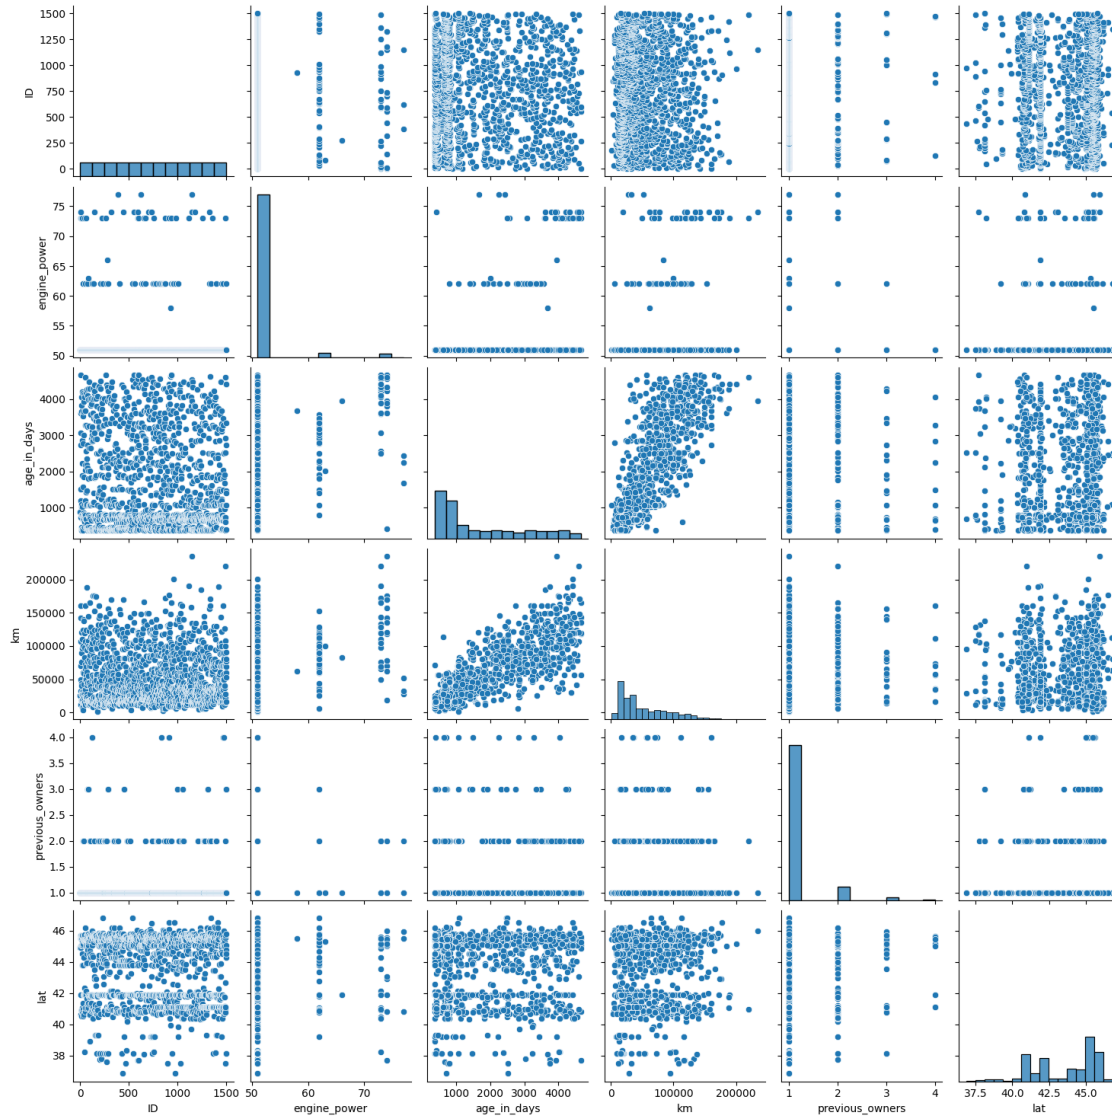
```
[11]:
```

```
[12]: df1=df1[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
         'lat']]
```

2 EDA AND VISUALIZATION

```
[13]: sns.pairplot(df1)
```

```
[13]: <seaborn.axisgrid.PairGrid at 0x7e6dfe587430>
```



```
[14]: sns.distplot(df1['km'])
```

<ipython-input-14-ad27032804f7>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df1['km'])
```

```
[14]: <Axes: xlabel='km', ylabel='Density'>
```



```
[15]: sns.heatmap(df1.corr())
```

```
[15]: <Axes: >
```



3 TO TRAIN THE MODEL AND MODEL BUILDING

```
[16]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   ID              1500 non-null   float64
1   engine_power    1500 non-null   float64
2   age_in_days     1500 non-null   float64
3   km              1500 non-null   float64
4   previous_owners 1500 non-null   float64
5   lat             1500 non-null   float64
dtypes: float64(6)
memory usage: 70.4 KB
```

```
[17]: x=df1[['ID', 'age_in_days', 'km','previous_owners',  
        'lat']]  
y=df1['engine_power']
```

```
[18]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
[19]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
[19]: LinearRegression()
```

```
[20]: lr.intercept_
```

```
[20]: 50.613595484526456
```


```
[21]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
[21]:
```

	Co-efficient
ID	-0.000267
age_in_days	0.000831
km	0.000004
previous_owners	-0.292565
lat	0.004989

```
[22]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
[22]: <matplotlib.collections.PathCollection at 0x7e6dfaecd8d0>
```

4 ACCURACY

```
[23]: lr.score(x_test,y_test)
```

```
[23]: 0.1174629752442442
```

```
[24]: lr.score(x_train,y_train)
```

```
[24]: 0.09327242421014503
```

```
[25]: from sklearn.linear_model import Ridge,Lasso  
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
[25]: Ridge(alpha=10)
```

```
[26]: rr.score(x_train,y_train)
```

```
[26]: 0.09326942619342626
```

```
[27]: rr.score(x_test,y_test)
```

[27]: 0.11747171879864649

```
[28]: la=Lasso(alpha=10)  
      la.fit(x_train,y_train)
```

[28]: Lasso(alpha=10)

```
[29]: la.score(x_train,y_train)
```

[29]: 0.09245387333030353

```
[30]: la.score(x_test,y_test)
```

[30]: 0.11694127378125141